

AMATH 582 Homework 3: Principal Component Analysis

Eric A. Silk

Abstract—Principal Component Analysis (PCA) is a highly useful technique for exploratory data analysis which decomposes a series of observations along axes of the most variance, which then allows for discarding of the less important variables to project into a lower dimensional subspace. This report explores its use in tracking the motion of paint can.

Index Terms—Singular Value Decomposition, Principal Component Analysis, Object Tracking, OpenCV

I. INTRODUCTION

PPrincipal Component analysis is a technique used to determine the best method of representing data in a lower dimensional subspace. It works by using an orthogonal transformation to convert data which may be correlated into a series of vectors which are uncorrelated. In a sense, it can produce the “best” basis by which to represent a set of data.

In order to explore this method, we will perform PCA upon the motion of a flashlight attached to a paintcan, moving in a variety of patterns (simple oscillation, oscillation and pendulum motion, oscillation, pendulum, and spinning), measured by cameras in three arbitrary positions and orientations.

II. THEORETICAL BACKGROUND

A. PCA

The technique of Principal Component Analysis works by computing the eigenvalues and eigenvectors of the covariance matrix of the data matrix \mathbf{X} , which is defined as $\mathbf{X}\mathbf{X}^T$. Calculating this directly can be numerically problematic – instead, the Singular Value Decomposition can be used.

B. SVD

For full details of the computation, see Dr. Kutz’s Book, “Data-Driven Modeling and Scientific Computation.”

In short, however, SVD of a data matrix produces three matrices:

$$svd(\mathbf{X}) = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^* \quad (1)$$

These resulting matrices can be interpreted within the context of this exercise as follows:

- 1) \mathbf{U} : Each column contains the direction of the PCA modes.
- 2) $\mathbf{\Sigma}$: Singular values on the diagonal, representing the energy of each mode.
- 3) \mathbf{V} : Each column represents the displacement in time along a mode

In an example provided by Kelsey Maass on Piazza, the meaning of the \mathbf{U} and \mathbf{V} matrix are swapped. This is because the data matrix \mathbf{X} arranges the observations as columns, rather than rows, which is opposite the convention used in Dr. Kutz’s book.

C. Properties

One of the most important properties of this technique is a relative independence from the bases used in the initial observation of the data. For instance, in this homework, we are given three perspectives that are clearly non-orthogonal. Because of our a priori knowledge of the system, we could conceivably transform the data to force a better perspective. In many cases, this isn’t possible. Use of SVD/PCA will allow us to do this more generally. More importantly, PCA will tell us which of the modes are most critical in describing the phenomena. This can be used in a number of ways:

- 1) Data Exploration, to find the “important” modes to investigate further
- 2) Feature Engineering, to feed a more limited subset of data to a machine learning algorithm
- 3) Visualization, to project a high dimension dataset into a visualizable subspace
- 4) Lossy compression, by discarding the lower energy modes

III. ALGORITHM IMPLEMENTATION AND DEVELOPMENT

A. Object Tracking

Object tracking proved to be a large portion of the project. In the end, OpenCV with Python bindings was used to identify and track an object of interest.

1) *Thresholding*: In order to facilitate object isolation, and given that the paint can had a lit flashlight on top of it, a simple thresholding of the grayscale image was optionally applied. Specifically, a “to zero” threshold was used. This works by identifying regions of luminance less than the threshold and setting them to zero. Comparatively, a binary threshold sets areas below the threshold to zero and areas above to the maximum luminance.

2) *Bounding Box Selection*: The first frame of the video is held on screen until the user draws a bounding box around the object of interest and presses either the space or enter key. This works quite well in most cases, except in instances in which the primary point of interest (the lit side of the flashlight) is not immediately visible.

3) *Tracking*: Once the bounding box is selected, a tracker provided by OpenCV was used. Specifically, the CSRT tracker, or “Discriminative Filter with Channel and Spatial Reliability” tracker (I’m unsure how the acronym matches, but alas). This was the best choice in a rather non-rigorous series of tests, and is touted as a fairly accurate method (albeit more computationally expensive than some other objects).

One major issue that was noted was, in the event of object occlusion or other loss of tracking, there was very poor recovery behavior. If this was to be in production, some other method of object detection would be required, or manual re-selection of the object upon failure.

4) *Saving of Data*: Upon conversion completion, any failures to track would be recorded as NaN, and the total number of these were reported. If the user felt they were acceptable, the data was then saved as .npz file for easy import.

B. SVD/PCA

1) *SVD*: Fortunately, per instructor permission, the SVD itself was able to be done using a pre-built implementation. Specifically, `numpy.linalg.svd`, which wraps LAPACK’s `_gesdd` routine, allowing for high numerical precision and speed.

The observation vectors were stacked before calculating the SVD, truncating to the length of the shortest for a given case.

2) *Plotting and Exploration*: Following the example set in Dr. Kutz’s book, a series of four plots were made for each case:

- 1) Modal Energy
- 2) Modal Energy, on SemiLog axes
- 3) Modal behavior of the 2 most significant modes
- 4) Temporal behavior of the 2 most significant modes

Additionally, reconstruction of the SVD with only N modes was performed to observe the effects of discarding lower energy modes.

IV. COMPUTATIONAL RESULTS

A. Kutz Plots

The results as “Kutz Plots” (described in “Plotting and Exploration”) for each of the cases can be seen in Figures 1, 3, 4, and 6.

The first results (Fig. 1) appear fairly straightforward. We know that only one major mode exists in the data – strong motion in the Z direction; however, PCA indicates there are two modes. Why is this occurring? Little to no effort was made to synchronize the videos during image tracking. As such, looking at the temporal behaviors of each mode, we can observe what appears to be a “lag” or “phase shift”. I believe this is due to the same phenomena (the oscillation) being observed at a slightly delayed time.

It is also worth noting that the modal directions themselves are comprised primarily of one or two directions. This corresponds to the known orientation of the cameras and the motion relative to that.

If the videos were perfectly synchronized, through the use of a “clapper” or similar, I’m certain these two modes would

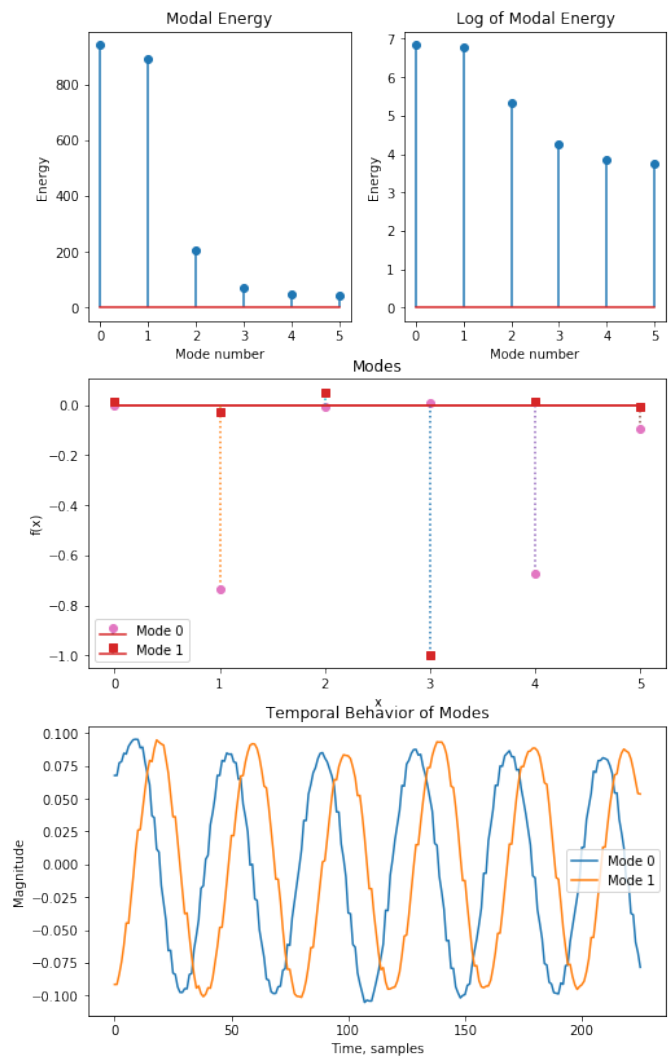


Fig. 1. Kutz Plot, Case 1

collapse into one strongly dominant mode. Alternatively, if time permitted, manual re-alignment in post could achieve a similar goal.

This is supported by inspecting Fig. 2 and noting that while camera 1’s y location and camera 3’s x location both start in a “trough” whereas camera 2’s y location begins at a “peak”, suggesting out of phase behavior.

Based upon this, the results in the shaken camera case (Fig. 3) are reasonable. The introduction of noise raises the modal energy floor. Instead of the ideal of one clearly dominant mode (or even 2, as explained in case 1), secondary modes also become candidates. Additionally, the modal directions become more mixed, with contributions from more modes. This makes sense, as the introduction of shake looks like travel along other dimensions.

Introducing another dimension of motion in the form of pendulum-like oscillations, along with the previously stated issues of synchronization, produce results that aren’t unexpected. Ideally, two modes would dominate in Fig. 4, but instead 3 seem to be of significance. Again, I believe this can be explained by comparing the relative peaks and valleys

Case 1




Fig. 2. X and Y positions of the object in each camera, Case 1

of the X and Y coordinates in each of the camera's recordings (Fig. 5).

Finally, case 4 introduces rotation. Since this is being mapped to rectangular coordinates, I would expect to see an additional mode corresponding to the "depth" direction. Looking at the resulting modal directions and temporal behavior of each component, this seems correct. The added modal directions of each mode are largely independent of each other. The first two temporal trends appear to be oscillatory, corresponding to the spring and the pendulum action, while the third shows a "flattening" at several points (most notably in the span between samples 100 and 150, likely corresponding to the paint can's rotation reversing).

V. SUMMARY AND CONCLUSION

From these

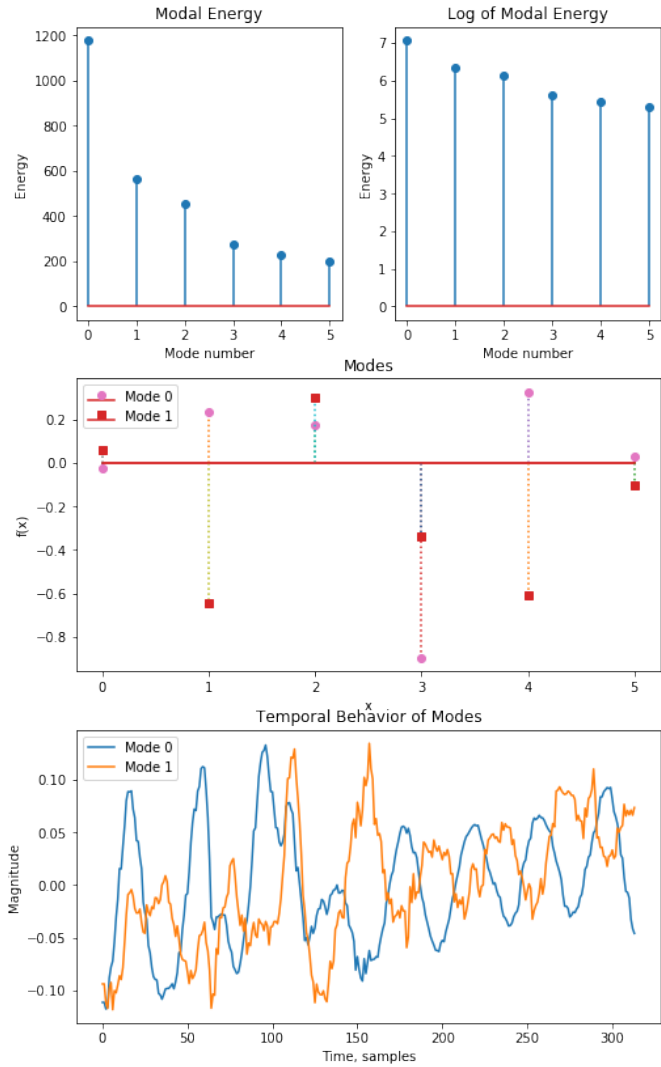


Fig. 3. Kutz Plot, Case 2

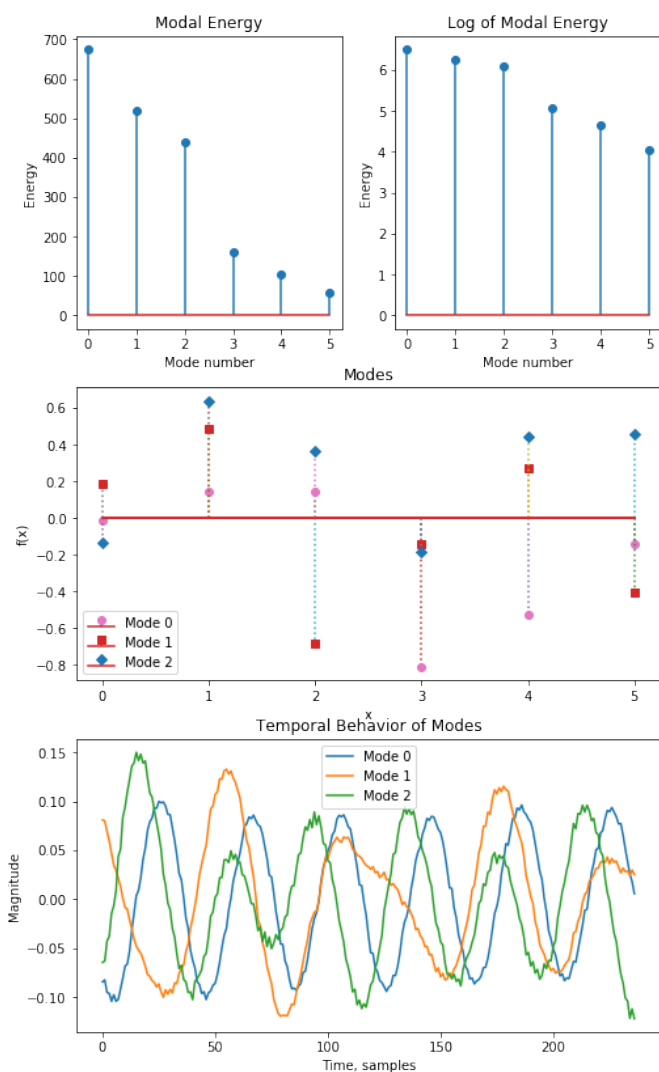


Fig. 4. Kutz Plot, Case 3




Fig. 5. X and Y positions of the object in each camera, Case 3

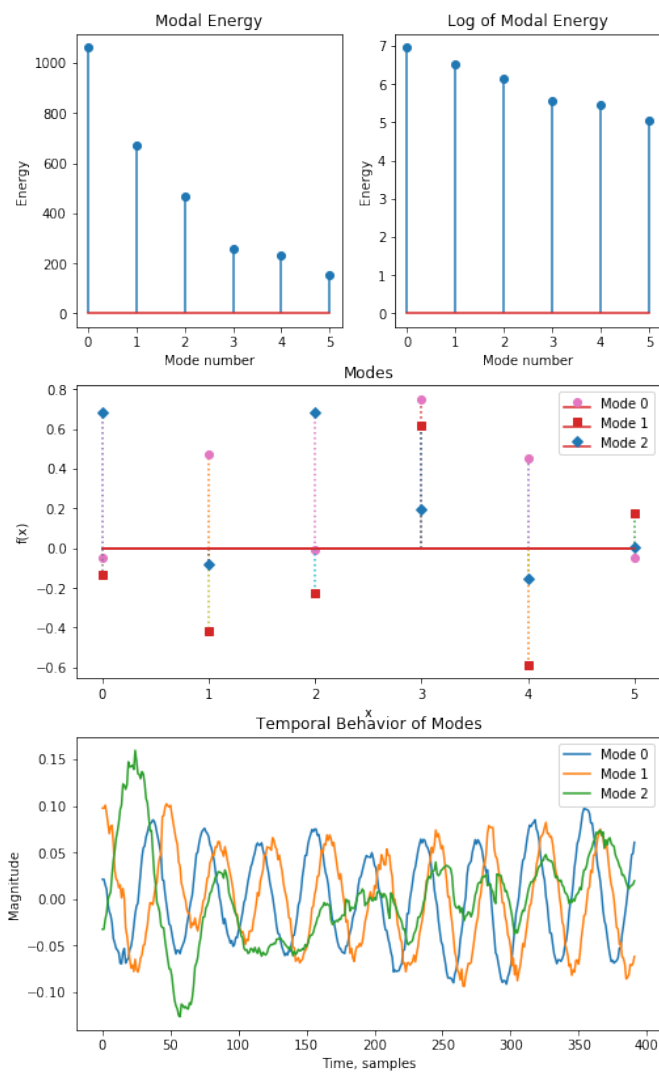


Fig. 6. Kutz Plot, Case 4

VI. APPENDIX A: FUNCTIONS USED

Commonly used functions, such as plotting, routine mathematical manipulations, and others, are not included here as they are deemed trivially understood.

A. *scipy.io.loadmat*

Given a path to a *.mat file, load the contents as a dictionary. Includes the header and global variables that may be included.

B. *cv2.cvtColor*

Converts an image or video to a new color representation, such as RGB, BGR, or Grayscale.

C. *cv2.threshold*

Given a grayscale image or video and a luminance threshold, apply a thresholding function. Options include binary, to zero, and others.

D. *cv2.imshow*

Show an image using CV2.

E. *cv2.selectROI*

Assuming an image is displayed, allow the user to select a bounding box "region of interest".

F. *cv2.TrackerCSRT_create*

Create a CSRT tracker. Must be initialized with a bounding shape and initial frame.

G. *cv2.rectangle*

Draw a rectangle on the active frame with the given corner coordinates, color specification, and line thickness.

H. *cv2.waitKey*

Waits the given time in milliseconds. If a key is pressed, its value is returned. Used for checking for early quits ("q" key, here).

I. *cv2.destroyAllWindows*

Method to assert the dominance of *nix systems over weak, MSDOS systems. Kidding, really its an active method to remove all open windows managed by OpenCV.

J. *numpy.load*

Loads a *.npy file. Conversions were done on the original *.mat files to make them easier to import.

K. *numpy.squeeze*

Given an ndarray with any axes of size 1, return a copy of the data with the redundant axes removed. E.g., a (3,1) matrix will become (3,).

L. *numpy.stack*

Given a sequence of arrays in tuple form, join them along a given axis. Used to combine the individual observations into a data matrix.

M. *numpy.linalg.svd*

Compute the Singular Value Decomposition of the input matrix, and returns the U , Σ , and V^* matrices.

VII. APPENDIX B: PYTHON CODE

See my [Github](#) for the full repository, including this source code for this IEEE template \LaTeX document. Code is also attached at the end of this report.